



Department of Electrical Engineering and Computer Science

## MASSACHUSETTS INSTITUTE OF TECHNOLOGY

### 6.033 Computer Systems Engineering: Spring 2011

# Quiz 3

There are 15 questions and 10 pages in this quiz booklet. Answer each question according to the instructions given. You have **90 minutes** to answer the questions.

Some questions are harder than others and some questions earn more points than others—you may want to skim all questions before starting.

For true/false questions, you will receive 0 points for no answer, and negative points for an incorrect answer. Do not guess; if you are unsure about your answer, consult your notes. The total score for each numbered question (1 through 15) will be at least 0 (i.e., those scores cannot go negative).

If you find a question ambiguous, be sure to write down any assumptions you make. **Be neat and legible.** If we can't understand your answer, we can't give you credit!

**Write your name in the space below.** Write your initials at the bottom of each page.

**THIS IS AN OPEN BOOK, OPEN NOTES, OPEN LAPTOP QUIZ, BUT  
DON'T USE YOUR LAPTOP FOR COMMUNICATION WITH OTHERS.**

**CIRCLE your recitation section number:**

- 10:00** 1. Lampson/Pesterev
- 11:00** 2. Lampson/Pesterev 3. Ports/Mutiso
- 12:00** 4. Ports/Mutiso
- 1:00** 5. Katabi/Raza 6. Strauss/Narula
- 2:00** 7. Katabi/Raza 8. Strauss/Narula

*Do not write in the boxes below*

1-4 (xx/24)	5-7 (xx/19)	8-10 (xx/18)	11-12 (xx/12)	13-15 (xx/27)	Total (xx/100)

**Name:**

**I Reading Questions**

1. [8 points]: Based on the paper “The Recovery Manager of the System R Database Manager”, which of the following statements are correct?

(Circle True or False for each choice.)

- A. **True / False** As part of recovery from a crash, System R restores the current version of all shadowed files to their shadow copy.
- B. **True / False** The shadow copy of a file reflects only the effects of committed transactions.
- C. **True / False** System R will undo the effects of certain transactions during recovery.
- D. **True / False** System R will redo the effects of certain transactions during recovery.

2. [6 points]: Based on the paper “Beyond Stack Smashing: Recent Advances in Exploiting Buffer Overruns”, which of the following statements are correct?

(Circle True or False for each choice.)

- A. **True / False** The basic form of attack described in the paper (stack smashing) allows an attacker to execute arbitrary code by overwriting far enough beyond the end of an application buffer to also overwrite the stack pointer.
- B. **True / False** Even if a system guards against buffer overruns by making the stack non-executable, and uses stack cookies to guard against arc injection / return-into-libc attacks, the system could still be vulnerable to attacks which modify function pointers.
- C. **True / False** A heap buffer overrun writes beyond the end of a dynamically-allocated object (e.g., created with `new` in C++ or `malloc()` in C), instead of a procedure’s local variables.

**Initials:**

**3. [6 points]:** Which of the following statements is true about the Porcupine distributed email system described by Saito, Bershad and Levy?

**(Circle True or False for each choice.)**

- A. True / False** In Porcupine, soft state such as the user map is replicated on every node for fault-tolerance.
- B. True / False** In order to balance the load across the cluster, when a new message arrives Porcupine always picks a lightly loaded node at random and adds the message to the user's mailbox fragment on that node.

**4. [4 points]:** Porcupine's scheme for maintaining cluster membership chooses a new coordinator whenever anyone detects an event: a node going up or down. The coordinator broadcasts a ping and gets a response from every node that is up. The total amount of work done in a cluster of  $n$  nodes, each time some node goes up or down, is proportional to which of the following?

**(Circle the BEST answer)**

- $n$
- $n \log n$
- $n^2$

**Initials:**

## II Lectures

**5. [4 points]:** According to Prof. Abelson's lecture, the "Good Samaritan Provision" of the Communications Decency Act provided protection to which of the following parties?

**(Circle True or False for each choice.)**

- A. True / False** CMU so that it can study the use of pornographic information on the Internet without being liable.
- B. True / False** AOL in the case against Ken Zeran, because AOL merely distributed information but was not responsible for the content.
- C. True / False** Prodigy in the case versus Stratton-Oakmont because Prodigy filtered offensive content.
- D. True / False** The nastiest place on earth because it merely distributed content that others generated.

**6. [8 points]:** Consider the mechanisms for storing and checking user's passwords, as described in lecture.

**(Circle True or False for each choice.)**

- A. True / False** Rather than storing user's login passwords directly, Unix stores the result of applying a hash function to the user's password. One reason for this approach is so that users' plaintext passwords are not revealed if an attacker manages to read the password database.
- B. True / False** Unix combines *salts* with a user's password before sending them through the hash function to enact a policy that requires users to periodically revise their passwords.
- C. True / False** Let  $H()$  be a hash function. If given a value  $p$ , it is easy to find a value  $P$  such that  $H(P)$  equals  $p$ , then  $H()$  is a poor choice of hash function to use as part of a password checking system.
- D. True / False** The study of password use described in lecture found that different people almost never choose the same passwords.

**7. [7 points]:** Consider the Same Origin Policy (SOP) used for browser security.

**(Circle True or False for each choice.)**

- A. True / False** Javascript code from the page at `http://web.mit.edu/6.033` can access the DOM of another page at `http://web.mit.edu/6.041` loaded in the same browser. Assume the second page runs no Javascript code.
- B. True / False** Javascript code from the page at `http://web.mit.edu/6.033` can access the DOM of another page at `http://courses.csail.mit.edu/6.033` loaded in the same browser. Assume the second page runs no Javascript code.

**Initials:**

- C. True / False** Javascript code from the page at `http://web.mit.edu/6.033` can access the DOM of another page at `http://www.google.com` loaded in the same browser. Assume the second page runs no Javascript code.
- D. True / False** Suppose the 6.033 staff copies the Javascript code for an advertisement from `attacker.com` into the page at `http://web.mit.edu/6.033` (for example, by including a tag like `<SCRIPT SRC="http://attacker.com/ad.js">`). The Same-Origin Policy will allow that Javascript code from `attacker.com` to access the DOM of the 6.033 page (`http://web.mit.edu/6.033`).
- E. True / False** The Same-Origin Policy allows Javascript code from `site1.com` to access the cookie of `site2.com`, as long as the two sites are concurrently loaded in two different tabs in the same window.

**Initials:**

### III Worm Trouble

It is May 2011 and there is a new worm, *nitty*, released on the Internet. Security experts quickly analyze the code corresponding to Nitty to be as follows:

```
1 int X = 0; /* X is a global variable */
2
3 int rand() {
4     /* Note that 32-bit integers obviate the need for a modulus operation here. */
5     /* A = 214013, B = 2531011. */
6     X = X * A + B;
7     return X;
8 }
9
10 srand(seed) {
11     X = seed;
12 }
13
14 main() {
15     srand(get_tick_count());
16     while (1) { /* infinite loop */
17         for (int i = 0; i < 20,000; ++i) {
18             dest_ip = rand();
19             dest_port = rand() [0:15]; /* low 16 bits */
20             packetsize = 1000; /* 1000 bytes */
21             packetcontents = top_of_stack;
22             sendto();
23         }
24         int rand_temp = rand();
25         if (open(physicaldisk, rand_temp[13:15]))
26             overwrite_block(rand_temp[0:14] || 0x4e20);
27     }
28 }
```

You learned about network telescopes in 6.033 and decided to see if you can apply the concepts to infer information about this new worm. You set up a router at MIT and configure it to listen on two unused portions of the MIT IP address space corresponding to 18.32.0.0/11 and 18.96.0.0/11. This worm appears to differ from the Witty worm in lines 16 (always loop back to the same place), 18, 20 (always choose the same packet size), and 24–26.

**8. [4 points]:** What fraction of the total 32 bit IP address space does your network telescope listen on?

**Initials:**

**9. [8 points]:** Say your network telescope receives two packets from an infected node with the following parameters: (dest\_ip\_1, dest\_port\_1) and (dest\_ip\_2, dest\_port\_2) respectively. Given the two received packets, how would you identify whether they are likely to be from a single for loop execution?

**10. [6 points]:** Finally, you would like to measure the bandwidth of the infected node. To do this you look at two packets, from a single for loop execution, you received at your network telescope at times  $T_1$  and  $T_2$ . The two packets had the following parameters: (dest\_ip\_1, dest\_port\_1) and (dest\_ip\_2, dest\_port\_2). Assume the packet header size is 100 bytes, how would you measure the bandwidth of the infected node?

**Initials:**

## IV Two-Phase Commit

Consider the two-phase commit protocol for executing atomic transactions across multiple sites. This protocol is described in Sections 9.6.2 and 9.6.3 of the text, as well as in lecture, and is summarized below

One node is designated as the *coordinator*, and the other nodes are designated as *workers*. Under normal circumstances, the protocol proceeds as follows:

- The coordinator sends a PREPARE message to each worker.
- The worker responds by writing a PREPARED record to its log, then sending a VOTE COMMIT message to the coordinator.
- After receiving VOTE COMMIT responses from all workers, the coordinator writes a COMMITTED record to disk and sends a COMMIT message to each worker.
- When a worker receives this message, it writes a COMMITTED record to its log.

**11. [8 points]:** Which of the following statements about this protocol are true? Assume that none of the components (network, coordinator, and workers) crash permanently.

**(Circle True or False for each choice.)**

- A. True / False** If any worker commits a transaction (by writing COMMITTED to its log), all other workers will also eventually commit that transaction.
- B. True / False** It guarantees that all workers will commit the transaction within a fixed time.
- C. True / False** If a worker is using two-phase locking to guarantee before-or-after atomicity, it cannot release the locks held by a transaction until after it receives the COMMIT message from the coordinator.
- D. True / False** If the coordinator fails to receive a VOTE COMMIT message from some worker in a timely manner, it can decide to abort the transaction.

**12. [4 points]:** What is the earliest point after which the transaction is guaranteed to commit? Again, assume that none of the components (network, coordinator, and workers) crash permanently.

**(Circle the BEST answer)**

- After the coordinator sends PREPARE messages
- After the last worker sends its VOTE message
- After the coordinator writes the COMMITTED record to its log
- After all workers have written the COMMITTED record to their logs

**Initials:**



## V Consistency/Replication

The software that runs the retail web site for amazon.com uses an eventually consistent key-value storage system called Dynamo. The basic operations in such a system are to read and write the value for a particular key. For example, a customer's shopping cart is a single object, with the customer's user ID as its key and the set of items in the cart as the value. Versions are labeled by vector clocks, and an update overwrites any earlier versions with smaller clock labels. For simplicity, assume that a set operation contacts only one server, and that server then propagates the new value to the other replicas asynchronously. Each set operation can choose a different server (e.g., if it believes the last server it talked to may have failed or is otherwise unresponsive).

If nodes fail and recover, or respond slowly, it's possible to have two or more versions with clock labels that are not totally ordered; in this case it's the application's job to reconcile them when it reads the cart from Dynamo. In case of doubt the cart application reconciles by keeping any items whose status is in doubt.

**13. [8 points]:** Suppose the cart is empty and the user adds item A to the cart, then adds item B, then deletes A. What are the possible contents of the cart that a subsequent read could see?

**(Circle True or False for each choice.)**

- A. True / False    {A}
- B. True / False    {B}
- C. True / False    {A, B}
- D. True / False    {}

**14. [12 points]:** With the same sequence of operations, what values could a read see after the system reaches eventual consistency, and no more updates need to be propagated, if there are no other writes to the cart?

**(Circle True or False for each choice.)**

- A. True / False    {A}
- B. True / False    {B}
- C. True / False    {A, B}
- D. True / False    {}

**15. [7 points]:** It's possible to avoid these anomalies completely and have a strongly consistent (single-copy) storage system by running a replicated state machine, using a consensus protocol such as Paxos, to do the read and write operations. The disadvantages of this method, which presumably caused Amazon to reject it in favor of eventual consistency, include which of the following?

**(Circle True or False for each choice.)**

**Initials:**

- A. **True / False** You can't do writes without access to a majority of the replicas.
- B. **True / False** You can't do reads without access to a majority of the replicas.
- C. **True / False** It's more complicated for the application to deal with.
- D. **True / False** It requires the application to contact several servers directly.
- E. **True / False** The latency for both read and write operations is greater.

## End of Quiz

Please double check that you wrote your name on the front of the quiz,  
and circled your recitation section number.

**Initials:**