



Department of Electrical Engineering and Computer Science


MASSACHUSETTS INSTITUTE OF TECHNOLOGY

6.033 Computer Systems Engineering: Spring 2012

Quiz I Solutions

There are 13 questions and 9 pages in this quiz booklet. Answer each question according to the instructions given. You have **50 minutes** to answer the questions.

Grade distribution histogram:



I Reading Questions

1. [4 points]: Simon, in “The architecture of complexity,” claims that hierarchical systems evolve more quickly than non-hierarchical systems of comparable size. Which of the following arguments does he say support this claim.

(Circle True or False for each choice.)

A. **True / False** The observation that complex systems are more comprehensible if they are neatly decomposable.

Answer: False.

B. **True / False** The parable of the watchmakers.

Answer: True.

C. **True / False** Thermodynamic considerations about entropy.

Answer: False.

D. **True / False** The distinction between state descriptions and process descriptions.

Answer: False.

2. [5 points]: Answer the following questions based on the X Window System paper.

(Circle True or False for each choice.)

A. **True / False** When a Web browser runs on a Unix workstation and displays its pages on that workstation using X, the browser is the client of both the web server and the X server.

Answer: True.

B. **True / False** The window manager must be built into the core of the X server because it can operate on the windows of multiple clients.

Answer: False.

C. **True / False** The X server of the 1980's has a complex management infrastructure for color map management because display controllers of that era could not provide enough memory to store the color of every pixel in each image.

Answer: True.

D. **True / False** The X server informs its clients when a region of one of their windows becomes obscured so that the clients can stop sending update requests for that region.

Answer: False.

E. **True / False** Synchronization errors could happen between the X server and its clients when network latencies delayed client responses.

Answer: True. See section 9.3 in the paper.

Initials:

3. [6 points]: Answer the following questions based on the Unix paper.
(Circle True or False for each choice.)

- A. True / False** Checking the return value of the `fork()` function enables a child process to execute different instructions from its parent.

Answer: True.

- B. True / False** Since a child process can write to all files that are open by its parent at the time of the `fork`, these writes can create a race condition with writes from the parent.

Answer: True. The two processes can write to their respective file descriptors in any order, since there is no file locking in the original version of Unix.

- C. True / False** One of the advantages of multitasking is that it makes the system more responsive to user inputs.

Answer: True.

4. [6 points]: This question is in the context of the Eraser paper. Assume a multi-threaded program has three locks: `mu`, `mu0`, and `mu1`, as well as an array `a` with two locations, `a[0]` and `a[1]`. Whenever a thread is about to modify the whole array (i.e., both `a[0]` and `a[1]`) it acquires the lock `mu`, but whenever it is about to modify `a[0]` alone it acquires `mu0`, and whenever it is about to modify `a[1]` alone it acquires `mu1`.

- A. Yes / No** Could Eraser's lockset algorithm detect a race condition with respect to accesses to either `a[0]` or `a[1]`?

Answer: Yes. The lockset for `a[0]` will be empty, because in some cases it is accessed with `mu` held, and others with `mu0` held. Eraser's algorithm reports a possible race condition when the lockset becomes empty.

5. [6 points]: This question continues the Eraser question. Assume a multi-threaded program has one lock `mu2` and an array `a` with two locations, `a[0]` and `a[1]`. Whenever a thread is about to modify either `a[0]` or `a[1]` it acquires the same lock `mu2`.

- A. Yes / No** Could Eraser's lockset algorithm detect a race condition with respect to accesses to either `a[0]` or `a[1]`?

Answer: No. Eraser's lockset for both `a[0]` and `a[1]` will contain `mu2` and will never be empty for Eraser to report a race.

II Complexity

6. [10 points]: Chapter 1 of the text describes several techniques for coping with complexity: Modularity (M), Abstraction (A), Layering (L), Hierarchy (H), Design for iteration (D), and Indirection (I).

Initials:

For each of the following advantages, mark the appropriate letter (M, A, L, H, D, I) or N for “none of these” to say which technique *best* provides that advantage. For each question, there is only one best answer, but a given technique might be the best answer to more than one question.

- A. **M A L H D I N** Helps the designers incorporate feedback in system implementations.

Answer: D.

- B. **M A L H D I N** Helps simplify the task of debugging a complex system by letting implementers deal with smaller components

Answer: M.

- C. **M A L H D I N** Makes it easier for designers to take advantage of delayed binding in system implementations.

Answer: I.

- D. **M A L H D I N** Ensures that the implementation will obey the robustness principle.

Answer: A.

- E. **M A L H D I N** If this is done correctly, it can help reduce the number of inter-module interactions in large systems.

Answer: H.

III Names

7. [5 points]: One of the examples in the first hands-on exercise asked you to notice that even though your home directory might be `/mit/YOU`, the sequence `cd /mit/6.033; cd ../YOU` when executed in the `tcsh` shell does not get you to your home directory. However, if you perform the same experiment in `bash`, it works. From the viewpoint of our name resolution discussion, which statement is correct in `bash`?

(Circle the BEST answer)

- A. When executing the `cd ..` command, `bash` determines if any shorter symbolic links exist to the resulting directory and displays the shortest one.
- B. The `bash` shell uses not only the current directory as its name mapping context but also some additional history of how the current directory was reached.
- C. `Bash` uses only Unix pathnames, not inodes.
- D. None of the above.

Answer: B.

Initials:

IV Concurrency

In this question, you can assume that loads and stores to variables are atomic, and neither the compiler nor hardware will ever reorder instructions. `continue` transfers control flow to the beginning of the `while` loop.

Consider the following lock implementation using a variable `x` for threads numbered 1 through `n` where initially `x = 0`. Each thread's number is stored in variable `i`.

```
acquire():
    while True:
        if x != 0:
            continue ## retry
        x = i
        if x != i:
            continue ## retry
        return

release():
    x = 0
```

8. [10 points]: Which of the following is true for the above algorithm:
(Circle the BEST answer)

- A. It does not guarantee mutual exclusion.
- B. it guarantees mutual exclusion but not deadlock freedom.
- C. it guarantees both mutual exclusion and deadlock freedom.

Answer: A. Two threads can both check `x != 0` and succeed. Then, thread 1 will set `x = 1` and verify that `x != 1` is false. Then, thread 2 will set `x = 2` and verify that `x != 2` is false. Both then return from `acquire`.

Consider the following lock implementation using variables `x` and `y` for threads numbered 1 through `n`, where each thread's number is stored in variable `i`, and where initially `y = 0`:

```
acquire():
    while True:
        x = i
        if y != 0:
            continue ## retry
        y = 1
        if x != i:
            continue ## retry
        return

release():
    y = 0
```

9. [10 points]: Which of the following is true for the above algorithm:
(Circle the BEST answer)

Initials:

- A. It does not guarantee mutual exclusion.
- B. it guarantees mutual exclusion but not deadlock freedom.
- C. it guarantees both mutual exclusion and deadlock freedom.

Answer: B. The algorithm guarantees mutual exclusion but not deadlock freedom. On an intuitive level, the algorithm guarantees mutual exclusion because once some thread passes $y=1$, all threads that execute $x=i$ later will get stuck at $y!=0$, and among any collection of threads that pass the $y!=0$ test concurrently and race to get to $x=i$, only the one that wrote last in $x=i$ can get past $x=i$, and all later threads will never get past $y!=0$ until this winning thread leaves the critical section and executes release.

To give you a sense of what it takes to rigorously prove the correctness of this lock implementation, see a complete proof at http://web.mit.edu/6.033/2012/wwwdocs/assignments/s12_1_9.pdf.

V Client/Server and bounded buffers

Consider the following bounded buffer code (send and receive), assuming the variables `bb.in` and `bb.out` are 64 bits, never overflow, can be read and written atomically, and neither the compiler nor hardware will ever reorder instructions:

```

send(bb, m):
    // each invocation of send has
    // its own local variables:
    //   my_send_index (64-bit int)
    while True:
        acquire(bb.lock)
        if bb.in - bb.out < N:
            my_send_index = bb.in
            bb.in = bb.in + 1
            release(bb.lock)
            bb.buf[my_send_index mod N] = m
            return
        release(bb.lock)

receive(bb):
    // each invocation of receive
    // has its own local variables:
    //   my_rec_index (64-bit int)
    //   m (message)
    while True:
        acquire(bb.lock)
        if bb.in > bb.out:
            my_rec_index = bb.out
            bb.out = bb.out + 1
            release(bb.lock)

            acquire(bb.rec_lock)
            m = bb.buf[my_rec_index mod N]
            release(bb.rec_lock)

        return m
        release(bb.lock)

```

10. [12 points]: Which of the following is true for the above implementation:
(Circle True or False for each choice.)

- A. **True / False** The code is correct if there is one sender and one receiver executing at same time.

Initials:

Answer: False. This implementation of `send` and `receive` is broken in all cases. The `send` function increments `bb.in` and releases the lock without placing the message in the buffer, which means a concurrent `receive` can read an uninitialized message from the buffer at the `bb.in` location.

Since this code is not correct for a single sender and receiver, it is also incorrect for multiple senders or receivers.

B. True / False The code is correct if there is one sender and many receivers executing at same time.

Answer: False.

C. True / False The code is correct if there are many senders and one receiver executing at same time.

Answer: False.

D. True / False The code is correct if there are many senders and many receivers executing at same time.

Answer: False.

VI Operating systems

11. [8 points]: Circle all of the function calls that directly correspond to system calls in a Unix system (based on the Unix paper) in the following implementation of the cp program:

```
void cp(char *srcpath, char *dstpath) {
    int src = open(srcpath, O_RDONLY);
    int dst = open(dstpath, O_WRONLY);

    if (src < 0 || dst < 0)
        exit(-1);

    while (1) {
        char buf[1024];
        ssize_t cc = read(src, buf, sizeof(buf));
        if (cc <= 0)
            break;
        ssize_t n = write(dst, buf, cc);
        assert(cc == n);
    }

    close(dst);
    close(src);
}

int main(int argc, char **argv) {
    cp(argv[1], argv[2]);
    exit(0);
}
```

Answer: open, open, exit, read, write, close, close, exit.

Initials:

12. [10 points]: Suppose we run two user-mode programs, A and B, which are independent (e.g., which do not access any common files), on a Unix OS, and we run the Unix OS as a guest in a virtual machine (VM). What can fail if a bug (such as a divide-by-zero or a random memory write) appears in different components of the system? Assume there are no other bugs. Draw an X in the appropriate locations in the table below:

	Program A fails	Program B fails	Guest OS kernel fails	VMM fails
Bug in program A	X			
Bug in program B		X		
Bug in the guest OS kernel	X	X	X	
Bug in the VM monitor	X	X	X	X

Answer: see above.

13. [8 points]: Suppose that you discover a bug in the implementation of `read()` that is invoked by `cp`, and you want to fix this bug.

A. True / False Fixing this bug will require modifying the `cp` program.

Answer: False. The code for `read` is in the OS kernel.

B. True / False Fixing this bug will require modifying the OS kernel.

Answer: True. The code for `read` is in the OS kernel.

Initials: